



Face Attendance

Secure offline facial recognition & liveness detection for remote locations

- Fully offline face recognition + blink liveness (anti-spoofing)
- Sub-second on-device match · 6.5 MB model · React Native (Android / iOS)
- Offline-first attendance with sync & purge when connectivity returns

Model: w600k_mbf_float16.tflite (MobileFaceNet + ArcFace)

Validated 98.41% accuracy on held-out Indian faces · exceeds the >95% requirement


Problem & requirements

- Mark attendance at remote locations where connectivity is unreliable or absent.
- Recognition and anti-spoofing must run fully on-device, offline.
- Prevent fraud via photographs or screens (liveness / anti-spoofing).
- Lightweight model, fast (<1s) recognition, >95% accuracy on diverse Indian faces.
- Cross-platform React Native (Android + iOS); open-source components.
- Sync attendance to a server once connectivity returns, then purge local data.

Our approach in one line

An offline-first React Native app: on-device face embedding + blink liveness, a local outbox, and automatic upload-then-purge sync — proven end to end against a real REST endpoint.

System architecture



Everything above runs fully offline on the device.

After a match

Attendance (name, person id, timestamp) is written to a LOCAL outbox (AsyncStorage).
When the network returns, queued records are POSTed in batches to the REST endpoint.
On HTTP 200 the records are marked synced and PURGED from the device.

Recognition engine

- Model: w600k_mbf_float16.tflite — MobileFaceNet embedding network, ~6.5 MB.
- Input [1,112,112,3] RGB, normalized (pixel-127.5)/127.5; output 512-d embedding.
- Faces aligned to a fixed 5-point ArcFace template (umeyama similarity transform).
- Matching by cosine similarity of L2-normalized embeddings; configurable threshold.
- Same alignment + preprocessing shared by app and the offline benchmark pipeline.

Model contract

- Input: 112×112 RGB
- Normalize: $(x-127.5)/127.5$
- Output: 512-d, L2-normalized
- Match: cosine similarity
- Size: 6.5 MB (float16)
- Speed: <1s detect+align+embed

Liveness — blink detection

- Two-phase scan: lock onto a steady, centered face, then prompt the user to blink.
- A full open → closed → open eye transition is required to confirm a live person.
- Defeats photo and screen replay attacks — a static image cannot blink.
- Runs entirely offline using on-device face + eye-state detection.
- Liveness gate applies to attendance; enrollment uses the same guided scan without it.

Why blink (offline anti-spoofing)

The brief asks for offline anti-spoofing (e.g. blink, smile, head-turn). A required natural blink is robust, fast, and needs no network — rejecting printed photos and faces shown on a screen.

Performance & validation

6.5 MB

Model size

< 1 s

Recognition + liveness

98.41%

Accuracy (held-out)

97.62%

TAR @ FAR = 1%

Validation on held-out Indian faces

Evaluated the shipped model on a held-out Bollywood/Indian face set (identities unseen in any tuning).

Verification protocol: 101,592 genuine pairs and 20,000 impostor pairs across 34 held-out identities.

Best accuracy 98.41% · TAR@FAR=1% 97.62% · TAR@FAR=0.1% 93.52% — exceeds the >95% requirement.

A full fine-tuning pipeline was built and run; on this dataset it did not improve held-out accuracy, so the validated model is shipped. The pipeline is retained for larger-scale future adaptation.


Sync & purge (offline-first)

- Attendance is stored locally first — the app is fully usable with no connectivity.
- Connectivity is detected automatically (NetInfo); queued records upload in batches.
- REST contract: POST { records:[{id,personId,name,ts}, ...] }; server returns 200 OK.
- Idempotent by record id, so retries never double-count.
- On confirmation, synced records are PURGED from the device — no stale local data.
- Proven end to end against a real REST endpoint (uploaded N, purged N).

Production path

API Gateway → Lambda → DynamoDB over HTTPS; device uses a short-lived token / pre-signed URL — no AWS keys stored on the device.


Flow 1 — Enrolling a person



How enrollment works

A one-time setup per person: name + a guided multi-frame capture, averaged into a single on-device face profile. No blink is required for enrollment; profiles never leave the device.

Flow 2 — Marking attendance



Camera locks onto a steady, centered face (ring fills teal).


User blinks once — confirms a live person, defeats photos/screens.

Face matched on-device in under a second; name shown and queued.

How attendance works

Lock → blink liveness → on-device match, all offline. A successful match is written to the local outbox and uploaded automatically when the device is next online.

Flow 3 — Admin & sync



Enrollment, people and sync sit behind a PIN gate.

Enroll, manage people, view sync status, adjust settings.

Pending/synced counts; records upload then purge from the device.

How sync & purge works

Attendance is stored locally first; when connectivity returns it uploads in batches and, on server confirmation, is purged from the device — so no device retains attendance after upload.

Tech stack & status

- React Native 0.75 (Android 8+ / iOS 12+); camera via VisionCamera; detection via ML Kit.
- On-device inference via react-native-fast-tflite; AsyncStorage for local data.
- Connectivity via NetInfo; offline outbox → batched POST → purge on confirm.
- Offline training / benchmark pipeline (align, fine-tune, verify) in Python.

Status against the brief

Offline recognition (98.41% held-out) · Blink liveness (anti-spoof) · Sub-second compute
6.5 MB model · Sync & purge (proven against real endpoint) · Android + iOS source